



Introduction to Red Hat OpenShift Data Science


Cloud service for artificial intelligence
/machine learning (AI/ML) use cases

Eoin Crosbie
Technical Partner Enablement Manager

1. Why
2. What
3. Try

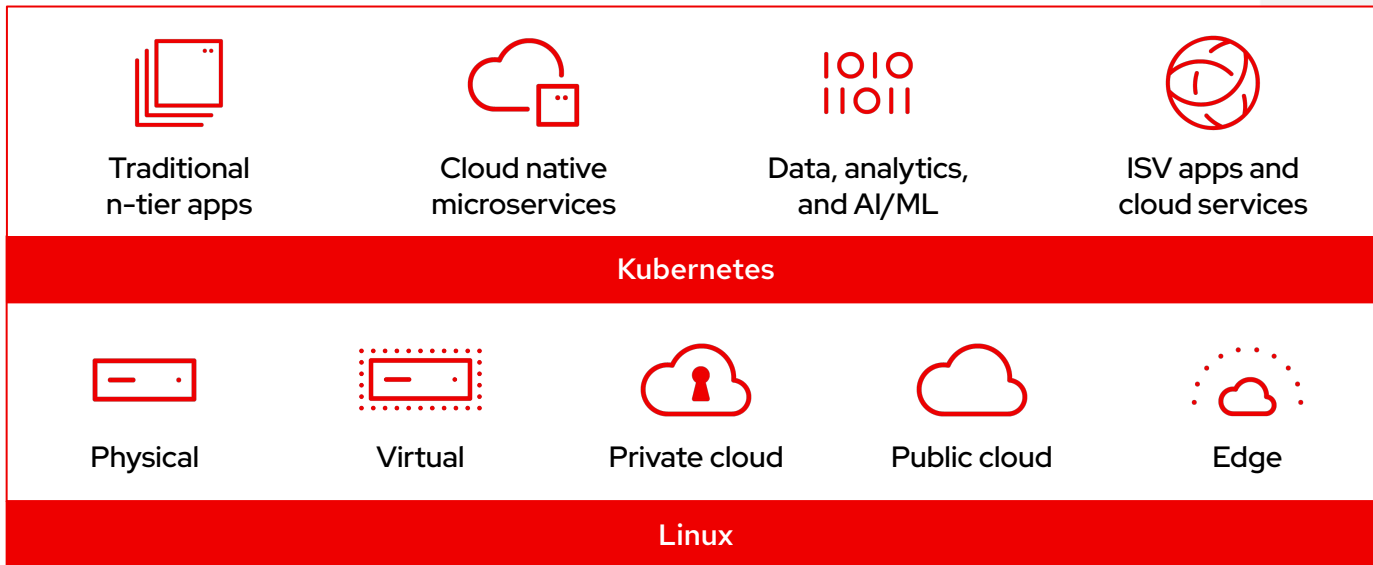
WHY





Open hybrid cloud

Red Hat helps customers build, run, and manage applications everywhere



Red Hat managed cloud services



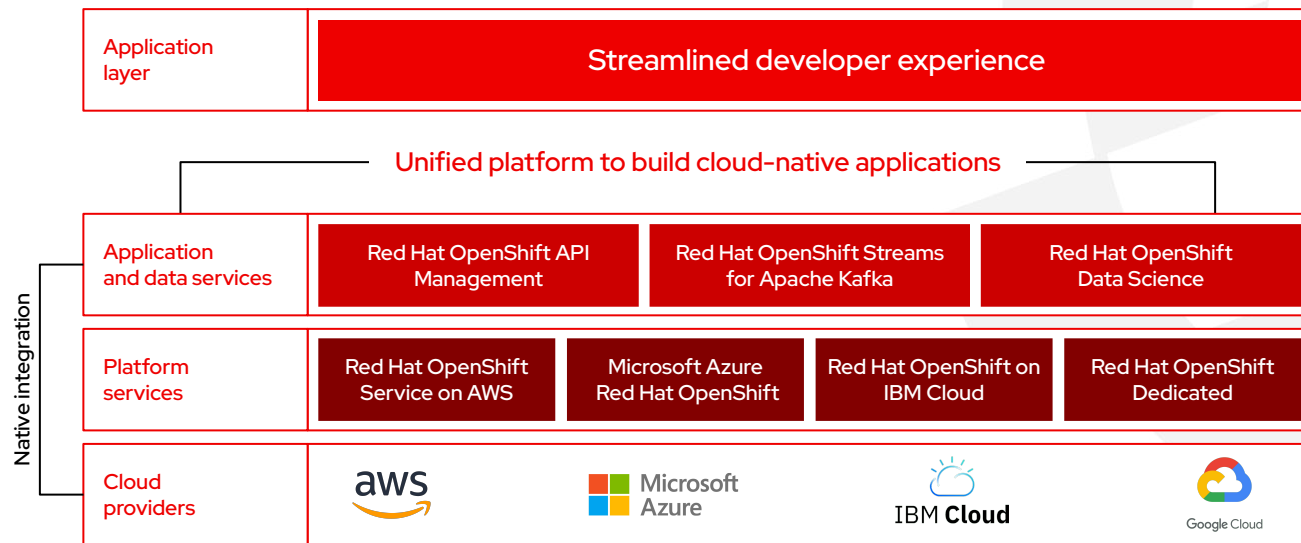
**Full stack management
and unified experience**




**Maximize full value of
Red Hat® OpenShift®**




Hybrid cloud flexibility



What came first,
The chicken or the E(d)gg?




Our approach to AI/ML



Hybrid cloud

Represents a workload requirement for our platforms across hybrid cloud.



Open source efficiency

Applicable to Red Hat's existing core business in order to increase open source development and production efficiency.



Intelligent platforms

Valuable as specific services and product capabilities, providing an intelligent platform experience.



Intelligent apps

Lets customers build intelligent apps using Red Hat products and our broader partner ecosystem.




Data as the foundation



Depth and scale without lock-in

Capabilities delivered through the combination of Red Hat and partner ecosystem




Managed cloud platform

Deployed on Red Hat OpenShift and managed on Amazon Web Services providing access to compute and accelerators based on your workload



Red Hat portfolio and services

Complement common data science tools in Red Hat OpenShift Data Science with other Red Hat products and cloud services

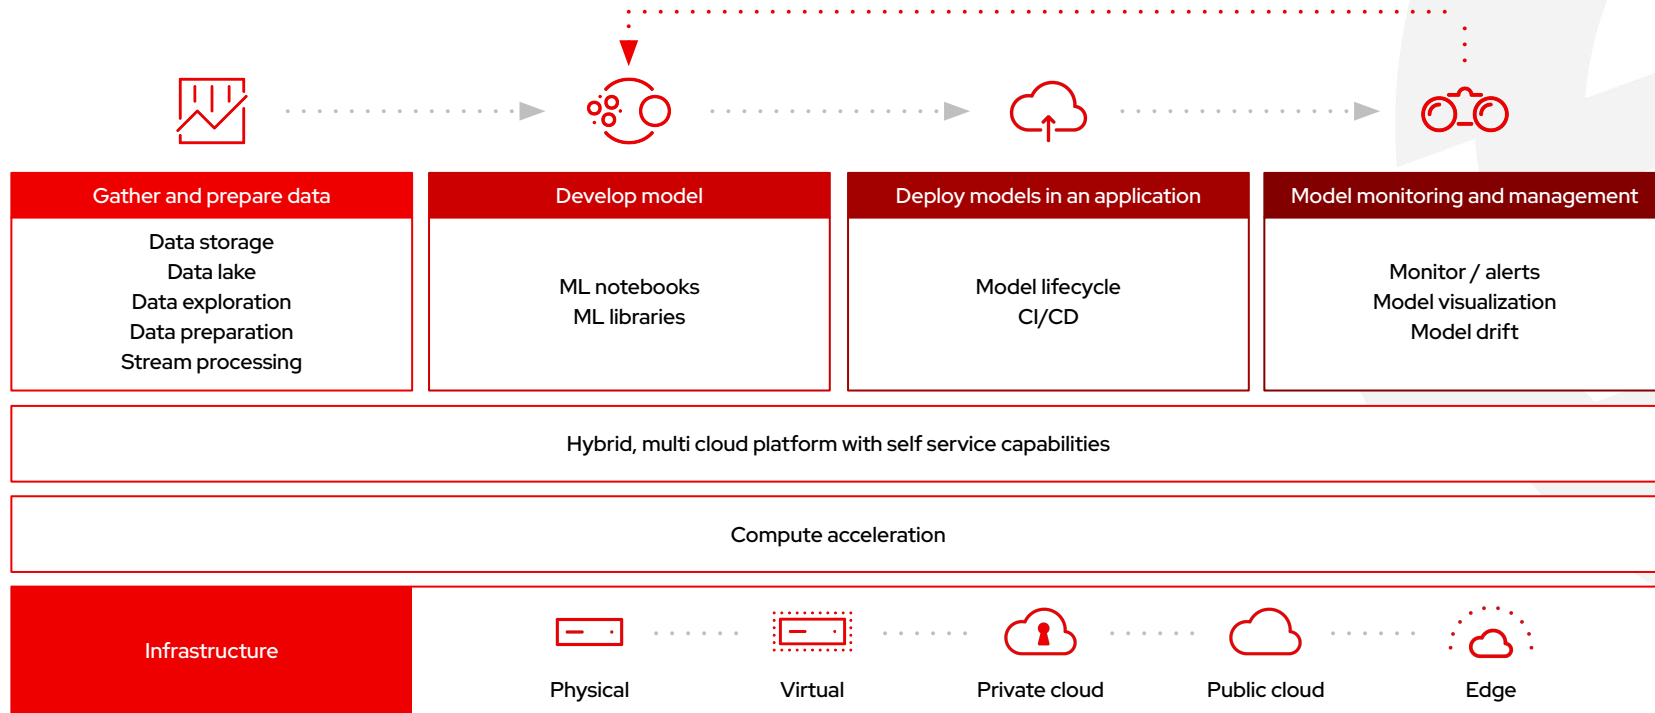


Partner ecosystem

Access specialized capabilities by adding certified ISV ecosystem products and services from Red Hat Marketplace

WHAT

Conceptual machine learning architecture



Key features of Red Hat OpenShift Data Science

Addressing AI/ML experimentation and integration use cases on a managed platform



Cloud Service

Available on Red Hat OpenShift Dedicated (AWS) and Red Hat OpenShift Service on AWS



Core data science workflow

Provides data scientists and intelligent application developers the ability to build, train, and deploy ML models



Increased capabilities/collaboration

Combines Red Hat components, open source software, and ISV certified software available on Red Hat Marketplace



Rapid experimentation use cases

Model outputs are hosted on the Red Hat OpenShift managed service or exported for integration into an intelligent application



Red Hat OpenShift Data Science

Tools and capabilities

Building on the foundations
of data science



Jupyter notebooks

Conduct exploratory data science in JupyterLab with access to core AI / ML libraries and frameworks including TensorFlow and PyTorch.



Source-to-image (S2I)

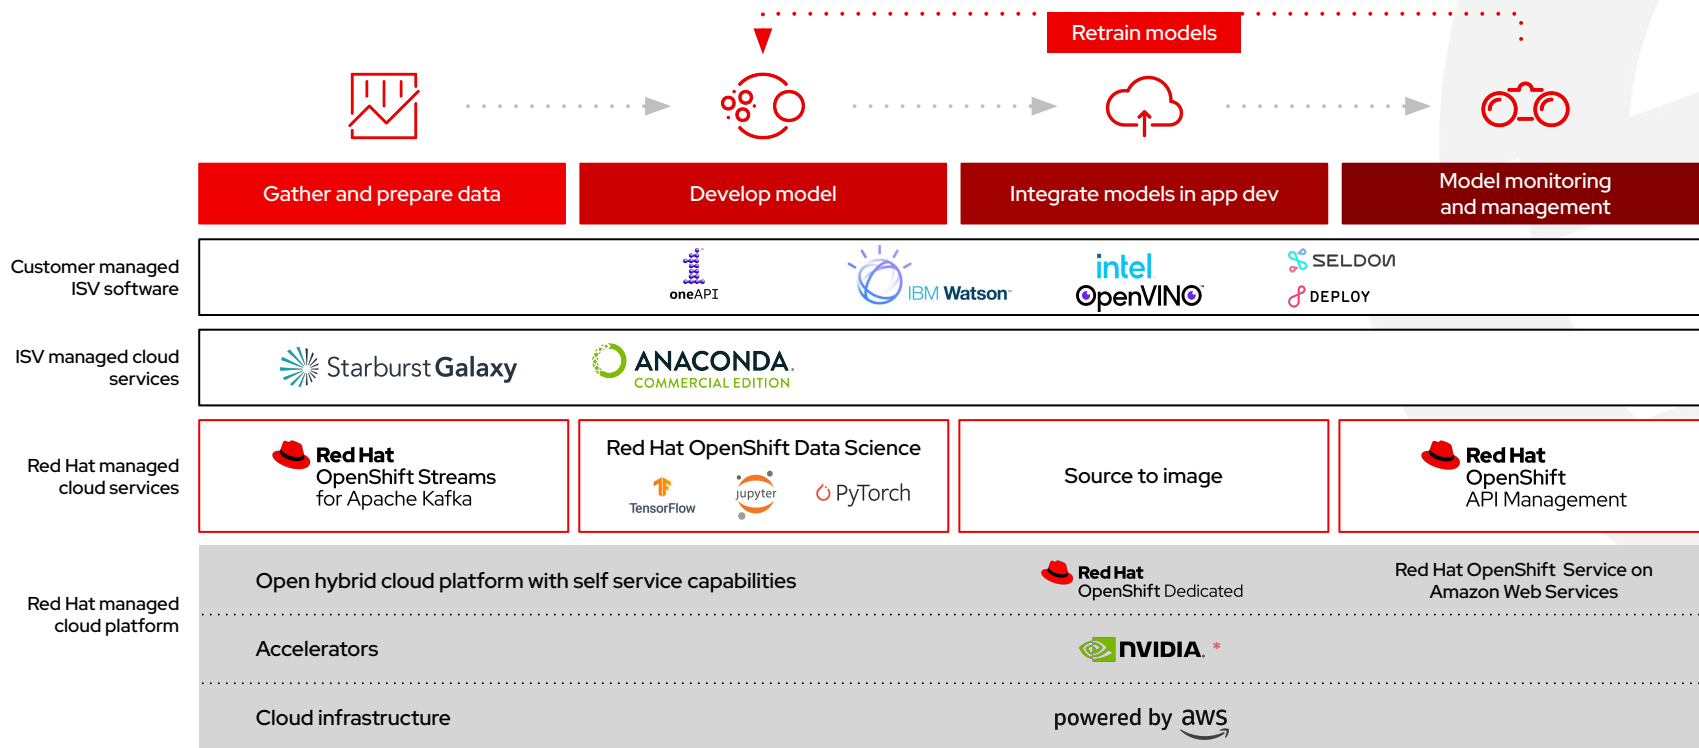
Publish models as end points via S2I for integration into intelligent apps. Rebuild and redeploy based on changes to the source code.



GPU Acceleration (post-initial release)

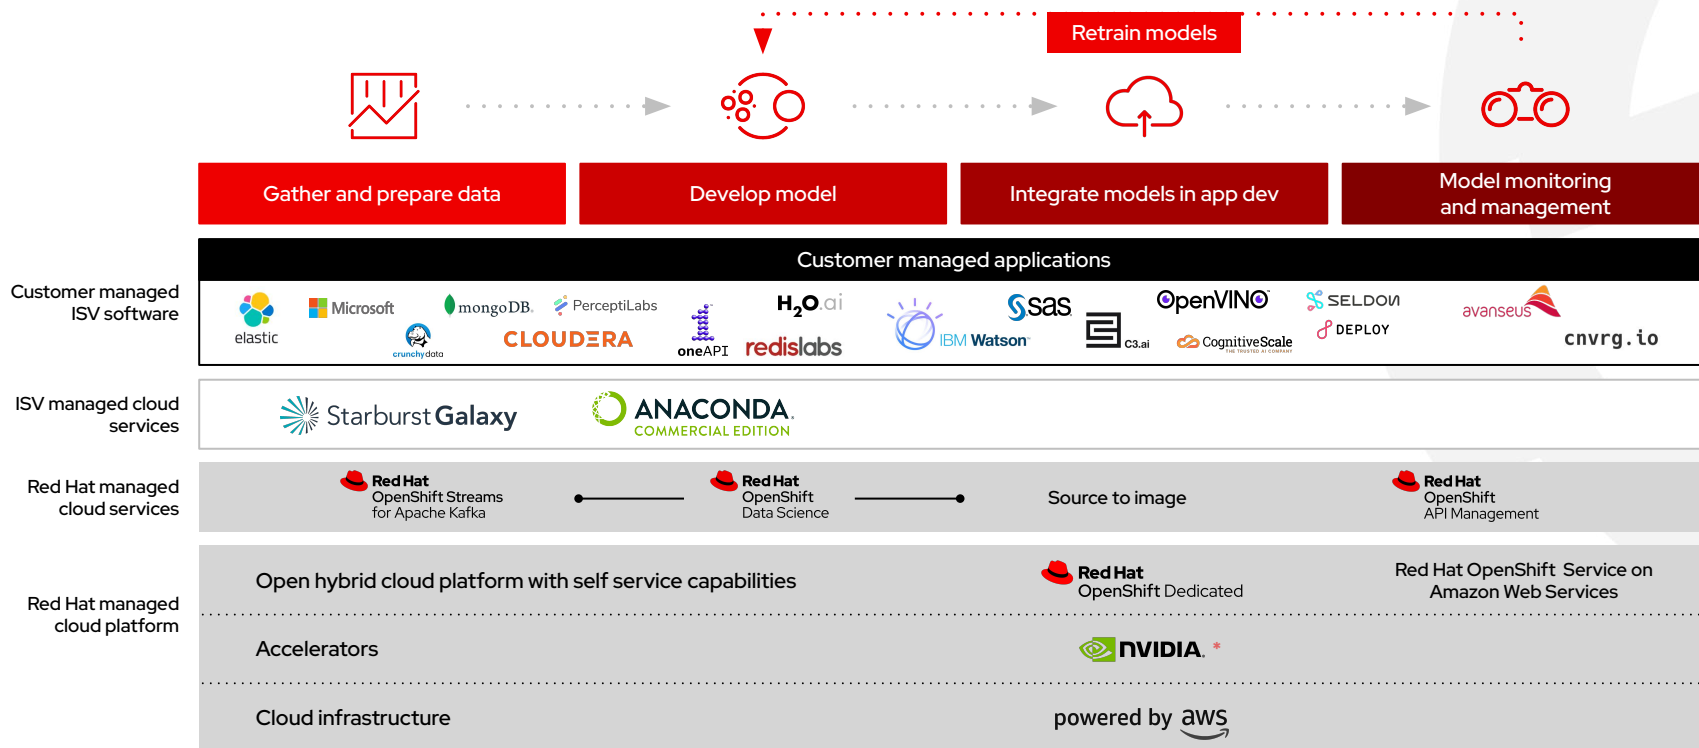
Accelerate your data science experiments through the use of GPU acceleration on the Red Hat OpenShift Dedicated platform.

Initial release components



*post initial release

... and integrating our partner ecosystem

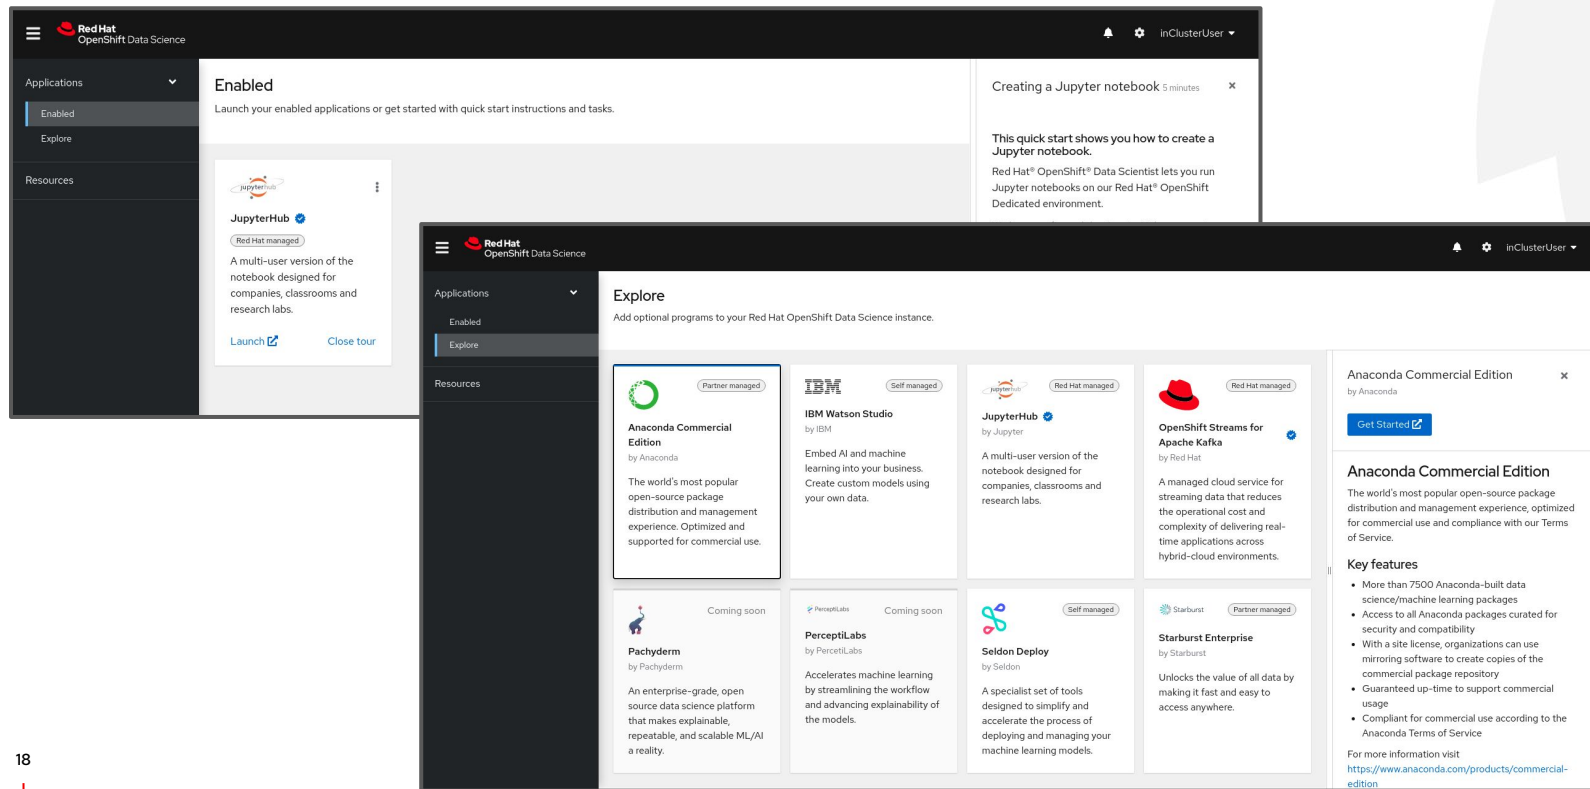


Dashboard resources

The screenshot shows the 'Resources' section of the Red Hat OpenShift Data Science dashboard. The left sidebar has 'Applications' and 'Resources' tabs, with 'Resources' selected. The main content area is titled 'Resources' and includes a search bar, sorting options ('Sort by name'), and filters ('Quick start X', 'Documentation X'). It displays 17 of 43 items in a grid of 10 cards. Each card features a logo, a title, a description, and a 'Start tour' link. The cards are arranged in two rows: the first row has 5 cards and the second row has 5 cards, with the last card in the second row being partially cut off.

Logo	Title	Description	Link
Anaconda	Anaconda Commercial Edition	The world's most popular open-source package distribution and management experience. Optimized and supported for commercial use.	Go to page
Red Hat	Connecting to Red Hat OpenShift Streams for Apache Kafka	This quick start will walk you through connecting to Red Hat Streams for Apache Kafka using Jupyter Notebooks	Start tour
JupyterHub	Creating a Jupyter notebook	This quick start will walk you through creating a Jupyter notebook.	Start tour
Anaconda	Creating an Anaconda-enabled Jupyter notebook	This quick start will walk you through creating an Anaconda-enabled Jupyter notebook.	Start tour
IBM	Deploying a Model with Watson Studio	This quick start walks you through importing a Notebook in Watson Studio, building a model with AutoAI, and deploying a model.	Start tour
JupyterHub	Deploying a sample Python application using Flask and OpenShift	How to deploy a Python model using Flask and OpenShift	
IBM	IBM Watson Studio	Embed AI and machine learning into your business. Create custom models using your own data.	
JupyterHub	JupyterHub	A multi-user version of the notebook designed for companies, classrooms and research labs.	
Scikit-Learn	Launch a SKLearn model and update model by canarying	How to perform a canary promotion of a Scikit-Learn model	
Scikit-Learn	Monitor drift for deployed model	Monitor drift for deployed image classifier model	

Dashboard user interface



The Red Hat Marketplace difference



Any cloud. On-premise. Anywhere OpenShift runs.

Build once, deploy to any environment with software that allows for workload portability across clouds.



Certified enterprise software

All software is certified for Red Hat OpenShift and uses Kubernetes operators for built-in management logic.



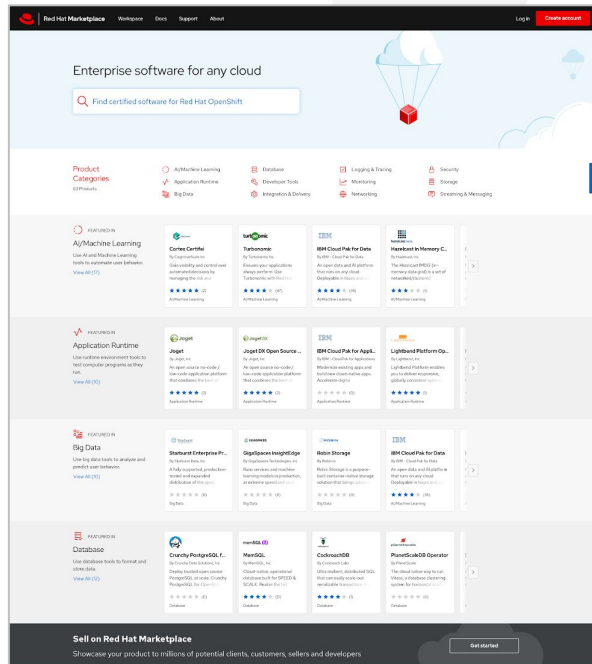
Continuous support

Central support portal offers a single point of entry for any software purchased in the marketplace.




Consolidated usage tracking & spend optimization

Pool spending across clouds. Monitor license usage, expiration, and renewals on a single dashboard.




How we differentiate versus alternatives



Hybrid cloud

Deploy models in containerized format for intelligent apps on-premise or in cloud




Fully managed cloud service

Management of entire Kubernetes stack, application services, and Jupyter-as-a-service



Open source innovation

Track changes and fixes to core open source AI/ML tooling and get access to upstream innovation



Partner ecosystem

Choice of over 30+ AI/ML technology partners as alternative to many monolithic frameworks

Red Hat OpenShift Data Science is HIPAA compliant

Ensure your cloud journey has the necessary privacy and security safeguards



Products that are HIPAA compliant are:

- ▶ Red Hat OpenShift Dedicated (AWS & GCP)
- ▶ Red Hat OpenShift Service on AWS
- ▶ Red Hat OpenShift API Management
- ▶ Red Hat OpenShift Streams for Apache Kafka
- ▶ Red Hat OpenShift Service Registry
- ▶ Red Hat OpenShift Data Science

Timeline

Beta

- ▶ Red Hat Summit announcement and demo
- ▶ Start Beta program with small group of customers
- ▶ Announce aligned consulting services

Future (Enhancements TBD)

- ▶ Support other cloud platforms
- ▶ Custom notebook images
- ▶ Additional AI / ML ISVs
- ▶ Enhance MLOps: KServe, Kubeflow Pipelines



April 2021

Early customer adoption

2H 2021

TBD

First release: Version 1

- ▶ As a cloud service add-on to OpenShift Dedicated (on AWS initially) or Red Hat OpenShift Service on AWS
- ▶ Initial AI / ML ISV partners
- ▶ Available via Red Hat Developer Sandbox

TRY

- Dev Sandbox

<https://developers.redhat.com/products/red-hat-openshift-data-science/overview>

- License Plate Detection

<https://redhat-scholars.github.io/rhods-lp-workshop/rhods-lp-workshop/index.html>




Red Hat OpenShift Data Science

Partner
Enablement Event
31st May



Red Hat OpenShift Data Science



Learn more ►



Contact us ►

Thank you

Red Hat is the world's leading provider of
enterprise open source software solutions.
Award-winning support, training, and consulting
services make
Red Hat a trusted adviser to the Fortune 500.

 linkedin.com/company/red-hat

 youtube.com/user/RedHatVideos

 facebook.com/redhatinc

 twitter.com/RedHat

The surrounding **infrastructure platform** and **processes** to support ML models are complex

